



UiT The Arctic University of Norway


Project 1: Boot-Up Mechanism

Loading a kernel and switching to 32-bit mode

Mike Murphy

UiT

Spring 2024



Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

Project 1

Project 1 is a demo

- ▶ You do not have to turn anything in
- ▶ Get the code. Build it. Run it.
- ▶ Set up your build environment
- ▶ Familiarize yourself with the code and the tools

We will be building an OS

- ▶ Each project will build on the previous
- ▶ Project 1: We give you a “hello world”
- ▶ Project 2: You start building a kernel
- ▶ Projects 3-6: You add more and more features

Project 1: “Hello World”

The code contains

- ▶ Bootloader
- ▶ “Hello world” kernel
- ▶ Utility to create a bootable disk image

The image

1. Boots
2. Loads the kernel
3. Prints a hello message

Architecture

- ▶ We’re going back to the 90s: 32-bit Intel x86 (aka i386)
- ▶ The OS can boot on some old hardware, but it’s getting harder
- ▶ We will develop primarily in an x86 emulator: Bochs

Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

What happens when a PC boots?

First steps

1. CPU: starts executing at 0xffffffff0 (*reset vector*)
2. Motherboard: makes sure there's something to execute: a jump to 0xf0000
3. BIOS ROM: starts at 0xf0000

BIOS: Basic Input Output System

- ▶ Firmware in ROM
- ▶ Knows how to talk to disks, keyboard, display
- ▶ Knows how to look for a bootable disk
- ▶ Provides interface for software to talk to disks, keyboard, display

Next steps

4. BIOS: looks for a bootable disk
5. BIOS: loads bootloader from disk
6. Bootloader: loads kernel from disk
7. Kernel: starts

UEFI is the New BIOS

UEFI: Unified Extensible Firmware Interface

- ▶ Began around 1998
- ▶ Replacement for BIOS
- ▶ Much more sophisticated

For example, the bootloader:

- ▶ BIOS:
 - ▶ Reads first 512 bytes of disk
 - ▶ Those 512 bytes contain simple partition table and bootloader code
 - ▶ First bootloader typically needs to load a *second stage* bootloader that understands the filesystems and can find the kernel
- ▶ UEFI:
 - ▶ Understands filesystems and executable file formats
 - ▶ Loads bootloader like an OS would load a normal program

We Don't Use UEFI

UEFI BIOS Compatibility

- ▶ UEFI replaced BIOS, but it kept *Compatibility Mode Support* (CSM) ...until ~2020
- ▶ New computers are UEFI-only

Our OS still uses the old BIOS boot system

- ▶ This is why it's getting harder to boot our OS on real PCs
- ▶ This is why we're giving you this code instead of making you write it
- ▶ Writing a BIOS bootloader requires knowledge of obsolete hardware modes

Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

x86 Family History

Year	CPU	Regs	Addresses	Hottest feature
1978	Intel 8086	16-bit	20-bit (1 MiB)	Segmented addressing
1982	Intel 286	16-bit	24-bit (16 MiB)	Memory protection
1985	Intel 386	32-bit	32-bit (4 GiB)	Virtual memory paging
...				
2003	AMD Opteron	64-bit	52-bit (4 PiB)	64 bits!

Note the shifting bottlenecks

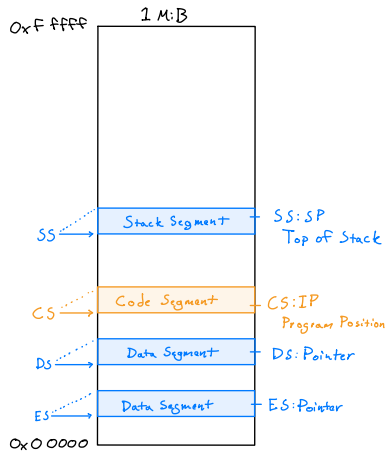
- ▶ 16-bit CPUs: registers. Harder to make wider registers than more memory
- ▶ 32-bit CPUs: memory, then regs. 4 GiB of RAM was a lot ... until it wasn't
- ▶ 64-bit CPUs: memory again. 4 PiB of RAM is a lot

Why is History Important?

Because every boot is a walk through history

- ▶ x86 CPUs are aggressively backwards compatible
- ▶ Even new CPUs act like an 8086 at boot, just in case
- ▶ Software must start with 16-bit code and bootstrap its way to 32- or 64-bit
 - ▶ BIOS-style: bootloader and OS must do this (our OS does this)
 - ▶ UEFI-style: firmware does this before handing off
- ▶ Understanding history helps you understand the present

i8086 Segmented Memory Addressing



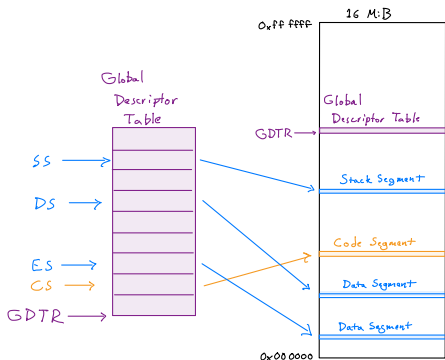
Segmented addressing

Address Space
20-bit (1 MiB)

Registers
16-bit (64 KiB)

- ▶ Divide memory into 64 KiB *segments*
 - ▶ Instruction pointer (IP) → Code Segment (CS)
 - ▶ Stack pointer (SP) → Stack Segment (SS)
 - ▶ Data pointers → Data Segments (DS, ES, FS, GS)
- ▶ CS, SS, DS, ES, FS, GS are 16-bit segment registers
 - ▶ Specify base of segment (shift by 4)
 - ▶ CS 0x1234 → Address 0x12340
 - ▶ IP 0x5678 → Address 0x12340 + 0x5678 = 0x179B8

i286 Protected Mode Addressing




Protected mode addressing

Address Space	Registers
24-bit (16 MiB)	16-bit (64 KiB)

- ▶ Add a layer of indirection
 - ▶ Move segment information into memory
 - ▶ Set up a table of 8-byte *segment descriptors*
- ▶ 8-byte segment descriptor in memory
 - ▶ 24-bit base address
 - ▶ 16-bit limit (end of segment)
 - ▶ various flags, including privilege level
- ▶ 16-bit segment register
 - ▶ 13 bit index into descriptor table
 - ▶ 3 bits for flags

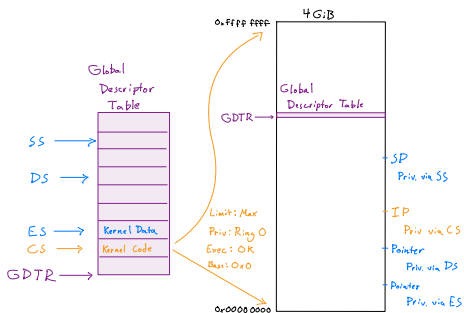
Protected Mode Privilege Levels



Privilege levels

- ▶ 2 bits for privilege level → 4 privilege levels
 - ▶ Ring 0 (innermost): operating system kernel
 - ▶ Rings 1 & 2: intended for device drivers
 - ▶ Ring 3 (outermost): user applications
- ▶ Segment descriptor determines current priv level
 - ▶ CS register selects descriptor
 - ▶ Descriptor priv level determines current priv level
 - ▶ Certain operations can only be executed with priv 0
 - ▶ Attempts to change CS are checked against current priv level vs requested priv level
- ▶ In practice, most operating systems use only kernel and user

i386 Flat Memory Model



Flat addressing

Address Space	Registers
32-bit (4 GiB)	32-bit (4 GiB)

► Do we still need segment descriptors?

1. To reach whole address space? — Not needed
2. To mark off protected regions of memory? — No, paging is better
3. To set current privilege level? — Still used

► Vestigial descriptor table

1. Kernel-level code: priv 0, executable
2. Kernel-level data: priv 0, no-execute
3. User-level code: priv 3, executable
4. User-level data: priv 0, no-execute

x86-64 Canonical Addressing

Address Space	Registers
52-bit (4 PiB)	64-bit (16 EiB)

- ▶ Pointers can address more memory than we can make RAM for

- ▶ CPU designs don't bother having 64 address lines
- ▶ Only 52 lines, which can address up to 4 PiB

- ▶ Canonical address form

- ▶ High bits must be all 0 or all 1

```
0xFFFF0 0000 0000 0000 -- 0xFFFF FFFF FFFF FFFF  
0x0000 0000 0000 0000 -- 0x000F FFFF FFFF FFFF
```

- ▶ This prevents software from trying to use unused bits for other purposes

Moving Through History on Boot

- ▶ CPU starts in 16-bit *Real Mode* (8086-style)
- ▶ To get to 32-bit *Protected Mode* (386-style)
 1. Create a Global Descriptor Table (GDT) in memory
 2. Set the Global Descriptor Table Register (GDTR)
 3. Enable the A20 line
 4. Enable Protected Mode
 5. Jump into 32-bit code

→ Our OS gets to here ←

- ▶ To get to 64-bit *Long Mode* (x86-64)
 1. Set up paging data structures in memory
 2. Set the appropriate paging registers
 3. Enable Long Mode
 4. Jump into 64-bit code

→ UEFI gets here before handing off ←

Side Note: AMD and x86-64

Why was it AMD who had the first 64-bit x86 system?

Because they developed the architecture

Intel Itanium and Intel Architecture 64 (IA-64)

- ▶ All-new architecture, not backwards compatible
- ▶ Based on *Very Long Instruction Word (VLIW)* concept
- ▶ Backwards compatibility with x86 through emulation layer, very slow
- ▶ Commercial flop. Tech press called it “The Itanic”.
- ▶ AMD extended the x86 architecture instead. That’s what stuck.
- ▶ Side-side note: UEFI was developed for Itanium

Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

BIOS Memory Map

- ▶ 1 MiB, divided into sixteen 64 KiB blocks

Addr	Block	Desc
0xf0000	F block	System ROM BIOS
0xe0000	E block	PCjr cartridges
0xd0000	D block	PCjr cartridges
0xc0000	C block	ROM expansion
0xb0000	B block	CGA memory
0xa0000	A block	EGA memory
0x00000	0-9 block	"conventional memory" / "low memory"

- ▶ Blocks 0-9: 640 KiB, this is the origin of the classic 640k limit


Where to Load Our OS



P1 memory map

- ▶ 0xa0000: Reserved: High memory
 - ▶ 0xb8000: VGA text buffer
- ▶ 0x80000: Stack
- ▶ 0x07c00: Boot block
- ▶ 0x01000: Kernel
- ▶ 0x00000: Reserved: Interrupt vector table

Disk Image: Normal OS for BIOS-Based System



Typical OS disk


Disk

- ▶ First 512 bytes: bootblock code + partition table
- ▶ Rest of disk is divided into partitions
- ▶ Partitions have filesystems
- ▶ Executables are ELF files

Boot

1. BIOS loads bootblock
2. Bootloader loads second-stage bootloader
3. Second-stage bootloader loads kernel ELF

Disk Image: Our OS



Our OS disk image

Disk

- ▶ First 512 bytes: bootloader code (no partition table)
- ▶ Rest of disk is raw kernel data

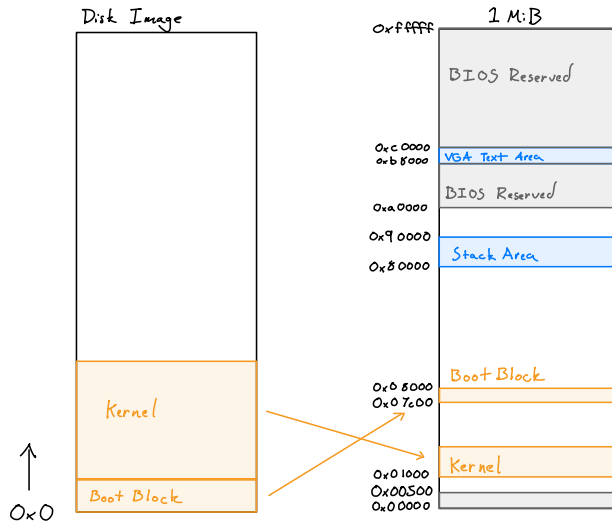
Boot

1. BIOS loads bootloader
2. Bootloader loads kernel

createimage program


- ▶ Read the kernel ELF file at build time
- ▶ Copy to image instead of memory

Disk Image to Memory



Loading boot block and kernel from disk to memory

Reading a Disk, BIOS Style: CHS Addressing



Spinning disk platters

- ▶ CHS: Cylinder / Head / Sector
- ▶ Floppy disks
 - ▶ Move arm to select *track*
 - ▶ Select one of two read *heads*, one per side
 - ▶ Time read to rotation to select a *sector*
- ▶ Spinning hard disks
 - ▶ Stack of disk platters
 - ▶ Move arm to select *cylinder* of stacked tracks
 - ▶ Select head to select platter and side
- ▶ Mapping blocks \leftrightarrow CHS
 - ▶ Block 0 \leftrightarrow CHS 0/0/1
 - ▶ No circuitry in floppy controller for mapping
 - ▶ Software must do the mapping
 - ▶ Even solid-state drives provide CHS emulation

Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

OS Source Code

Repository tree

```
.
|-- src
|   |-- boot    Bootloader source
|   |-- kernel  Kernel source
|   `-- lib     Libraries source
|-- host        Output directory for host machine (your computer)
|-- target      Output directory for target machine (x86-32)
`-- thirdparty  BIOS images for emulator
```

Build process

1. Compile `src/lib/**` `-> target/lib/lib*.a`
2. Compile `src/boot/**` `-> target/boot/bootblock`
3. Compile `src/kernel/**` `-> target/kernel/kernel`
4. Compile `src/createimage.c` `-> host/createimage`
5. Run `host/createimage` `-> image`

Make

Make Commands

```
make                # Default, equivalent to 'make image test'
docker-run make     # Run the same command in a Docker container
```

Makefiles

```
Makefile            # Main Makefile, delegates to host or target Makefiles
host/Makefile       # Rules for compiling for host machine
target/Makefile     # Rules for compiling for target machine
src/Makefile.common # Common rules, included by host and target Makefiles
```

Host System (Your Computer)

- ▶ x86-64 architecture
- ▶ Recommend Ubuntu 22.04 LTS (Jammy Jellyfish)
 - ▶ Other Linuxes should work, but may have GCC version issues
 - ▶ Windows Subsystem for Linux might work
 - ▶ Mac ???
- ▶ Can also use Docker
 - ▶ Dockerfile describes the system
 - ▶ `docker-run` script spins up a container and runs a command

```
docker-run make
```

Emulator: Bochs

- ▶ Supported emulator: Bochs
 - ▶ x86 emulator that focuses on accuracy over speed
 - ▶ Sorry, no QEMU (yet)
- ▶ Bochs built-in debugger
 - ▶ Similar to GDB, but a little clunky
 - ▶ Advantage: 16-bit code (bootloader)
 - ▶ Install `bochs` from `apt`: configured for built-in debugger
- ▶ Connect Bochs to GDB (recommended)
 - ▶ Much more powerful
 - ▶ Weakness: 16-bit code (because GDB assumes flat memory model)
 - ▶ Requires building Bochs from source

Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

Suggested Tasks

- ▶ Build: get code, set up environment, build OS
- ▶ Step through with a debugger
 - ▶ Bochs debugger? Step through bootloader (breakpoint: 0x7c00)
 - ▶ Bochs + GDB? Step through 32-bit code (breakpoint: `_start32`)
 - ▶ Get a feel for x86 ASM and how C maps to it
- ▶ Read code
 - ▶ Can you follow the bootblock ASM?
 - ▶ Can you follow the kernel C code?
 - ▶ Can you follow the C library code?
- ▶ Look at docs
 - ▶ `doc/x86`: x86 architecture and ASM programming
 - ▶ `doc/abi`: System V ABI docs (C calling convention, ELF format, etc.)
 - ▶ Download and flip through Intel manuals
 - ▶ Download and flip through AMD manuals (AMD may be easier to read)

Ask for Help

- ▶ The OS course is challenging
 - ▶ But it's supposed to be a fun challenge
- ▶ We're here to help
 - ▶ Ask questions on Discord (and please ask in the open chats)
 - ▶ Talk to your TA
 - ▶ Come to the colloquium sessions
- ▶ We also need feedback
 - ▶ The course is in transition, there will be hiccups
 - ▶ The code is in transition, there will be bugs
 - ▶ Let us know what is working and what is not