



UiT The Arctic University of Norway


## Project 1: Boot-Up Mechanism

Loading a kernel and switching to 32-bit mode

Mike Murphy

UiT

Spring 2024



## Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

# Project 1

## Project 1 is a demo

- ▶ You do not have to turn anything in
- ▶ Get the code. Build it. Run it.
- ▶ Set up your build environment
- ▶ Familiarize yourself with the code and the tools

## We will be building an OS

- ▶ Each project will build on the previous
- ▶ Project 1: We give you a “hello world”
- ▶ Project 2: You start building a kernel
- ▶ Projects 3-6: You add more and more features

# Project 1: “Hello World”

## The code contains

- ▶ Bootloader
- ▶ “Hello world” kernel
- ▶ Utility to create a bootable disk image

## The image

1. Boots
2. Loads the kernel
3. Prints a hello message

## Architecture

- ▶ We’re going back to the 90s: 32-bit Intel x86 (aka i386)
- ▶ The OS can boot on some old hardware, but it’s getting harder
- ▶ We will develop primarily in an x86 emulator: Bochs

## Project 1

Project Overview

**Boot Basics**

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

# What happens when a PC boots?

## First steps

1. CPU: starts executing at 0xffffffff0 (*reset vector*)
2. Motherboard: makes sure there's something to execute: a jump to 0xf0000
3. BIOS ROM: starts at 0xf0000

## BIOS: Basic Input Output System

- ▶ Firmware in ROM
- ▶ Knows how to talk to disks, keyboard, display
- ▶ Knows how to look for a bootable disk
- ▶ Provides interface for software to talk to disks, keyboard, display

## Next steps

4. BIOS: looks for a bootable disk
5. BIOS: loads bootloader from disk
6. Bootloader: loads kernel from disk
7. Kernel: starts

# UEFI is the New BIOS

## UEFI: Unified Extensible Firmware Interface

- ▶ Began around 1998
- ▶ Replacement for BIOS
- ▶ Much more sophisticated

For example, the bootloader:

- ▶ BIOS:
  - ▶ Reads first 512 bytes of disk
  - ▶ Those 512 bytes contain simple partition table and bootloader code
  - ▶ First bootloader typically needs to load a *second stage* bootloader that understands the filesystems and can find the kernel
- ▶ UEFI:
  - ▶ Understands filesystems and executable file formats
  - ▶ Loads bootloader like an OS would load a normal program

# We Don't Use UEFI

## UEFI BIOS Compatibility

- ▶ UEFI replaced BIOS, but it kept *Compatibility Mode Support* (CSM) ...until ~2020
- ▶ New computers are UEFI-only

## Our OS still uses the old BIOS boot system

- ▶ This is why it's getting harder to boot our OS on real PCs
- ▶ This is why we're giving you this code instead of making you write it
- ▶ Writing a BIOS bootloader requires knowledge of obsolete hardware modes



## Project 1

Project Overview

Boot Basics

**Some History**

BIOS Booting and Our OS

Building the OS

Project Summary

# x86 Family History

Year	CPU	Regs	Addresses	Hottest feature
1978	Intel 8086	16-bit	20-bit (1 MiB)	Segmented addressing
1982	Intel 286	16-bit	24-bit (16 MiB)	Memory protection
1985	Intel 386	32-bit	32-bit (4 GiB)	Virtual memory paging
...				
2003	AMD Opteron	64-bit	52-bit (4 PiB)	64 bits!

Note the shifting bottlenecks

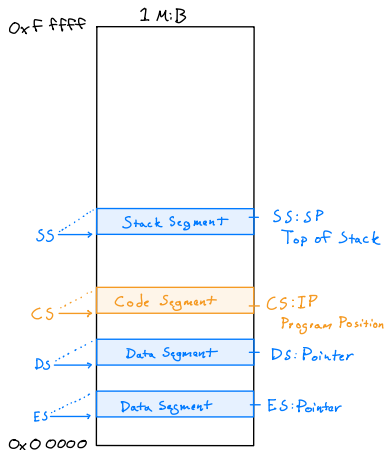
- ▶ 16-bit CPUs: registers. Harder to make wider registers than more memory
- ▶ 32-bit CPUs: memory, then regs. 4 GiB of RAM was a lot ... until it wasn't
- ▶ 64-bit CPUs: memory again. 4 PiB of RAM is a lot

# Why is History Important?

## Because every boot is a walk through history

- ▶ x86 CPUs are aggressively backwards compatible
- ▶ Even new CPUs act like an 8086 at boot, just in case
- ▶ Software must start with 16-bit code and bootstrap its way to 32- or 64-bit
  - ▶ BIOS-style: bootloader and OS must do this (our OS does this)
  - ▶ UEFI-style: firmware does this before handing off
- ▶ Understanding history helps you understand the present

# i8086 Segmented Memory Addressing

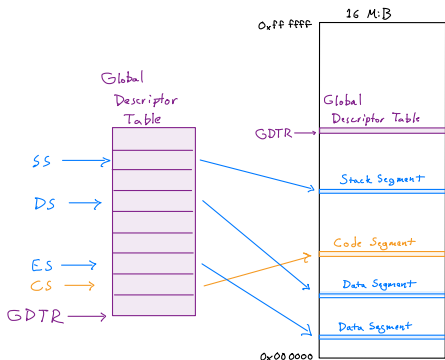


Segmented addressing

Address Space	Registers
20-bit (1 MiB)	16-bit (64 KiB)

- ▶ Divide memory into 64 KiB *segments*
  - ▶ Instruction pointer (IP) → Code Segment (CS)
  - ▶ Stack pointer (SP) → Stack Segment (SS)
  - ▶ Data pointers → Data Segments (DS, ES, FS, GS)
- ▶ CS, SS, DS, ES, FS, GS are 16-bit segment registers
  - ▶ Specify base of segment (shift by 4)
  - ▶ CS 0x1234 → Address 0x12340
  - ▶ IP 0x5678 → Address 0x12340 + 0x5678 = 0x179B8

# i286 Protected Mode Addressing




Protected mode addressing

Address Space	Registers
24-bit (16 MiB)	16-bit (64 KiB)

- ▶ Add a layer of indirection
  - ▶ Move segment information into memory
  - ▶ Set up a table of 8-byte *segment descriptors*
- ▶ 8-byte segment descriptor in memory
  - ▶ 24-bit base address
  - ▶ 16-bit limit (end of segment)
  - ▶ various flags, including privilege level
- ▶ 16-bit segment register
  - ▶ 13 bit index into descriptor table
  - ▶ 3 bits for flags

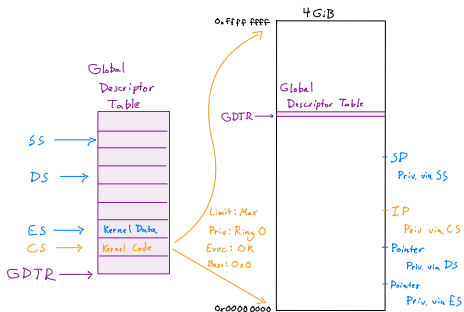
# Protected Mode Privilege Levels



Privilege levels

- ▶ 2 bits for privilege level → 4 privilege levels
  - ▶ Ring 0 (innermost): operating system kernel
  - ▶ Rings 1 & 2: intended for device drivers
  - ▶ Ring 3 (outermost): user applications
- ▶ Segment descriptor determines current priv level
  - ▶ CS register selects descriptor
  - ▶ Descriptor priv level determines current priv level
  - ▶ Certain operations can only be executed with priv 0
  - ▶ Attempts to change CS are checked against current priv level vs requested priv level
- ▶ In practice, most operating systems use only kernel and user

# i386 Flat Memory Model



Flat addressing

Address Space	Registers
32-bit (4 GiB)	32-bit (4 GiB)

## ► Do we still need segment descriptors?

1. To reach whole address space? — Not needed
2. To mark off protected regions of memory? — No, paging is better
3. To set current privilege level? — Still used

## ► Vestigial descriptor table

1. Kernel-level code: priv 0, executable
2. Kernel-level data: priv 0, no-execute
3. User-level code: priv 3, executable
4. User-level data: priv 0, no-execute

# x86-64 Canonical Addressing

Address Space	Registers
52-bit (4 PiB)	64-bit (16 EiB)

- ▶ Pointers can address more memory than we can make RAM for

- ▶ CPU designs don't bother having 64 address lines
- ▶ Only 52 lines, which can address up to 4 PiB

- ▶ Canonical address form

- ▶ High bits must be all 0 or all 1

```
0xFFFF0 0000 0000 0000 -- 0xFFFF FFFF FFFF FFFF  
0x0000 0000 0000 0000 -- 0x000F FFFF FFFF FFFF
```

- ▶ This prevents software from trying to use unused bits for other purposes



# Moving Through History on Boot

- ▶ CPU starts in 16-bit *Real Mode* (8086-style)
- ▶ To get to 32-bit *Protected Mode* (386-style)
  1. Create a Global Descriptor Table (GDT) in memory
  2. Set the Global Descriptor Table Register (GDTR)
  3. Enable the A20 line
  4. Enable Protected Mode
  5. Jump into 32-bit code

→ Our OS gets to here ←

- ▶ To get to 64-bit *Long Mode* (x86-64)
  1. Set up paging data structures in memory
  2. Set the appropriate paging registers
  3. Enable Long Mode
  4. Jump into 64-bit code

→ UEFI gets here before handing off ←

# Side Note: AMD and x86-64

Why was it AMD who had the first 64-bit x86 system?

Because they developed the architecture

## Intel Itanium and Intel Architecture 64 (IA-64)

- ▶ All-new architecture, not backwards compatible
- ▶ Based on *Very Long Instruction Word (VLIW)* concept
- ▶ Backwards compatibility with x86 through emulation layer, very slow
- ▶ Commercial flop. Tech press called it “The Itanic”.
- ▶ AMD extended the x86 architecture instead. That’s what stuck.
- ▶ Side-side note: UEFI was developed for Itanium

## Project 1

Project Overview

Boot Basics

Some History

**BIOS Booting and Our OS**

Building the OS

Project Summary


# BIOS Memory Map

- ▶ 1 MiB, divided into sixteen 64 KiB blocks

Addr	Block	Desc
0xf0000	F block	System ROM BIOS
0xe0000	E block	PCjr cartridges
0xd0000	D block	PCjr cartridges
0xc0000	C block	ROM expansion
0xb0000	B block	CGA memory
0xa0000	A block	EGA memory
0x00000	0-9 block	"conventional memory" / "low memory"

- ▶ Blocks 0-9: 640 KiB, this is the origin of the classic 640k limit


# Where to Load Our OS



P1 memory map

- ▶ 0xa0000: Reserved: High memory
  - ▶ 0xb8000: VGA text buffer
- ▶ 0x80000: Stack
- ▶ 0x07c00: Boot block
- ▶ 0x01000: Kernel
- ▶ 0x00000: Reserved: Interrupt vector table

# Disk Image: Normal OS for BIOS-Based System



Typical OS disk


## Disk

- ▶ First 512 bytes: bootblock code + partition table
- ▶ Rest of disk is divided into partitions
- ▶ Partitions have filesystems
- ▶ Executables are ELF files

## Boot

1. BIOS loads bootblock
2. Bootloader loads second-stage bootloader
3. Second-stage bootloader loads kernel ELF

# Disk Image: Our OS



Our OS disk image

## Disk

- ▶ First 512 bytes: bootloader code (no partition table)
- ▶ Rest of disk is raw kernel data

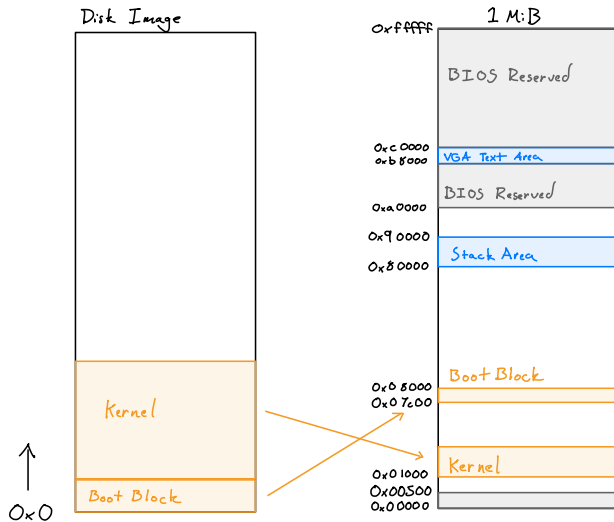
## Boot

1. BIOS loads bootloader
2. Bootloader loads kernel

## createimage program

- ▶ Read the kernel ELF file at build time
- ▶ Copy to image instead of memory


# Disk Image to Memory



Loading boot block and kernel from disk to memory



# Reading a Disk, BIOS Style: CHS Addressing



Spinning disk platters

- ▶ CHS: Cylinder / Head / Sector
- ▶ Floppy disks
  - ▶ Move arm to select *track*
  - ▶ Select one of two read *heads*, one per side
  - ▶ Time read to rotation to select a *sector*
- ▶ Spinning hard disks
  - ▶ Stack of disk platters
  - ▶ Move arm to select *cylinder* of stacked tracks
  - ▶ Select head to select platter and side
- ▶ Mapping blocks  $\leftrightarrow$  CHS
  - ▶ Block 0  $\leftrightarrow$  CHS 0/0/1
  - ▶ No circuitry in floppy controller for mapping
  - ▶ Software must do the mapping
  - ▶ Even solid-state drives provide CHS emulation

## Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

**Building the OS**

Project Summary

# OS Source Code

## Repository tree

```
.  
|-- src  
|   |-- boot    Bootloader source  
|   |-- kernel  Kernel source  
|   `-- lib     Libraries source  
|-- host        Output directory for host machine (your computer)  
|-- target      Output directory for target machine (x86-32)  
`-- thirdparty  BIOS images for emulator
```

## Build process

1. Compile `src/lib/**` -> `target/lib/lib*.a`
2. Compile `src/boot/**` -> `target/boot/bootblock`
3. Compile `src/kernel/**` -> `target/kernel/kernel`
4. Compile `src/createimage.c` -> `host/createimage`
5. Run `host/createimage` -> `image`

# Make

## Make Commands

<code>make</code>	# Default, equivalent to 'make image test'
<code>docker-run make</code>	# Run the same command in a Docker container

## Makefiles

<code>Makefile</code>	# Main Makefile, delegates to host or target Makefiles
<code>host/Makefile</code>	# Rules for compiling for host machine
<code>target/Makefile</code>	# Rules for compiling for target machine
<code>src/Makefile.common</code>	# Common rules, included by host and target Makefiles

# Host System (Your Computer)

- ▶ x86-64 architecture
- ▶ Recommend Ubuntu 22.04 LTS (Jammy Jellyfish)
  - ▶ Other Linuxes should work, but may have GCC version issues
  - ▶ Windows Subsystem for Linux might work
  - ▶ Mac ???
- ▶ Can also use Docker
  - ▶ Dockerfile describes the system
  - ▶ `docker-run` script spins up a container and runs a command

```
docker-run make
```

# Emulator: Bochs

- ▶ Supported emulator: Bochs
  - ▶ x86 emulator that focuses on accuracy over speed
  - ▶ Sorry, no QEMU (yet)
- ▶ Bochs built-in debugger
  - ▶ Similar to GDB, but a little clunky
  - ▶ Advantage: 16-bit code (bootloader)
  - ▶ Install `bochs` from `apt`: configured for built-in debugger
- ▶ Connect Bochs to GDB (recommended)
  - ▶ Much more powerful
  - ▶ Weakness: 16-bit code (because GDB assumes flat memory model)
  - ▶ Requires building Bochs from source

## Project 1

Project Overview

Boot Basics

Some History

BIOS Booting and Our OS

Building the OS

Project Summary

# Suggested Tasks

- ▶ Build: get code, set up environment, build OS
- ▶ Step through with a debugger
  - ▶ Bochs debugger? Step through bootloader (breakpoint: 0x7c00)
  - ▶ Bochs + GDB? Step through 32-bit code (breakpoint: `_start32`)
  - ▶ Get a feel for x86 ASM and how C maps to it
- ▶ Read code
  - ▶ Can you follow the bootblock ASM?
  - ▶ Can you follow the kernel C code?
  - ▶ Can you follow the C library code?
- ▶ Look at docs
  - ▶ `doc/x86`: x86 architecture and ASM programming
  - ▶ `doc/abi`: System V ABI docs (C calling convention, ELF format, etc.)
  - ▶ Download and flip through Intel manuals
  - ▶ Download and flip through AMD manuals (AMD may be easier to read)



# Ask for Help

- ▶ The OS course is challenging
  - ▶ But it's supposed to be a fun challenge
- ▶ We're here to help
  - ▶ Ask questions on Discord (and please ask in the open chats)
  - ▶ Talk to your TA
  - ▶ Come to the colloquium sessions
- ▶ We also need feedback
  - ▶ The course is in transition, there will be hiccups
  - ▶ The code is in transition, there will be bugs
  - ▶ Let us know what is working and what is not