

# Project 2: Cooperative Multitasking

Managing processes

INF-2201 Staff

Spring 2024

## Contents

<b>Project 2</b>	<b>2</b>
Project Overview . . . . .	2
Project 2: Administrative Info . . . . .	2
Project 2: Cooperative multitasking . . . . .	2
Project 2: Simple Environment . . . . .	2
Project Tasks . . . . .	3
Tasks Overview . . . . .	3
. . . . .	3
Processes and Threads / Process Control Block (PCB) . . . . .	3
Task: Design Process Control Block (PCB) . . . . .	3
Task: Initialize Kernel . . . . .	5
Memory Map . . . . .	5
Task: Implement Scheduler . . . . .	6
Task: Implement Context Switching . . . . .	6
Incorporating Assembly with C . . . . .	7
Task: Implement Locks: Blocking and Unblocking . . . . .	7
Task: Implement System Call Mechanism . . . . .	8
Syscall vs Scheduler . . . . .	8
Task: Measure Context Switch Time . . . . .	9
Extra Challenges . . . . .	9
Administrative Details . . . . .	9
Design Reviews . . . . .	9
Possible Topics for Design Review . . . . .	10
Code . . . . .	10
Report . . . . .	10
Hand In via Canvas . . . . .	10

# Project 2

## Project Overview

### Project 2: Administrative Info

- Mandatory assignment
  - This is a mandatory assignment
  - It will be graded as pass/fail by TAs
  - You must pass to gain admission to the exam
- Groups of two
  - You must work in groups of two
- Design review
  - Meet with TA and present your design
  - Informal, but must have some kind of slides / presentation
- Hand-in
  - Hand-in via Canvas
  - Code: whole repository (working tree + `.git/` dir)
  - Report: place in a `report/` directory in your repo
  - Zip up and upload to Canvas

### Project 2: Cooperative multitasking

- Switch between multiple threads of execution
- Threads can give up control via function:
  - `yield`: give up control
  - `block`: wait on a lock
  - `exit`: stop execution, allow kernel to clean up
- Kernel takes control and then gives it to another thread
  - Save context for yielding/blocking thread
  - Choose next thread
  - Set up / restore context for next thread

### Project 2: Simple Environment

- Protected mode, but no active protection
  - CPU in 32-bit Protected Mode
  - No protection active: all runs at kernel level (Ring 0)
  - Flat 32-bit address space
  - Kernel and processes share one address space
- No `malloc` / `free`
  - Statically allocate any globals / arrays you need
  - Allocate single structs from arrays




Figure 1: Returning to another thread

## Project Tasks

### Tasks Overview

- Design Process Control Block (PCB)
- Initialize kernel
- Implement scheduler
- Implement context switching
- Implement locks, blocking, and unblocking
- Implement system call mechanism
- Measure context switch time

### Processes and Threads / Process Control Block (PCB)

Process	Thread	File
Kernel	clock_thread	th1.c
Kernel	lock_thread[0-1]	th2.c
Kernel	mcpi_thread[0-3]	th3.c
Plane	main	process1.c
Math	main	process2.c

- `struct pcb` (Process Control Block)
  - One struct for process info + thread info
  - Don't let this blur your understanding of processes vs threads

### Task: Design Process Control Block (PCB)

#### We give you

- `struct pcb` in `pcb.h`
- next / previous pointers + fns to put PCBs into a linked list

```

struct pcb {
    /* PCBs as a doubly-linked list */
    struct pcb *next;
    struct pcb *previous;

    /* TODO */
};

```

```

Clock: 7 seconds
Lock test A: 200 OK
Lock test B: 200 OK

Pi A: 19/ 24->3.166667
Pi B: 20/ 24->3.333333
Pi C: 18/ 24->3.000000
Pi D: 21/ 24->3.500000
Pi T: 78/ 96->3.250000

Did you know: 1 + ... + 61 = 1891

== Process status ==
Pid Type St KStck
3 Thrd Run 23ee8
4 Thrd Run 24ee8
5 Thrd Run 25ee8
6 Thrd Run 26ee8
7 Proc Run 27ed0
8 Proc Run 29ed0

info : Kernel starting...
info : Initialized syscall entry fixed point.
info : Beginning task dispatch...

IPS: 215.146M | A: NUM | CAPS | SCRL | USB-HD |

```

Figure 2: Project 2 example screenshot




Figure 3: Project 2 processes and threads

- stack pointer? stack pointers?

### Look at other OSes

- What does Linux do? Windows?
- Read case studies in the textbook
- Search online
- Comparing OSes is a great thing to bring up in your design reviews

### Task: Initialize Kernel

- We give you:
  - `init_cpu()`: sets up Global Descriptor Table
  - `init_syscalls()`: sets up function pointer for system calls
- You must:
  - Initialize PCB table
  - Initialize PCBs for in-kernel threads and in-process threads
    - \* No need to do dynamic loading
    - \* Just set up an array at startup

### Memory Map




Figure 4: Project 1 memory map

- Kernel + Process 1 + Process 2 will overwrite boot block
- Bootblock must relocate itself
- This used to be a gotcha in the assignment

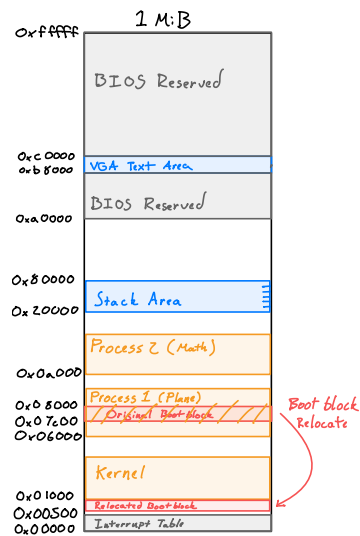


Figure 5: Project 2 memory map

- Stack area
  - Used for boot
  - Divide to allocate stacks to threads
  - How many stacks per thread?
- File: `addr.h`

### Task: Implement Scheduler

- How do you want to pick the next process?
- Simple round-robin is fine
- But you are welcome to try more advanced options

### Task: Implement Context Switching

- How do you switch between threads?
- Cooperative: thread gives up control
  - `yield`: give up control
  - `block`: wait on a lock
  - `exit`: stop execution, allow kernel to clean up
- What state to save?
  - Registers
  - Stack pointer
  - Anything else?
- Where and how to save it?

- Store in PCB fields?
- Push everything onto the stack?
- Hybrid?

## Incorporating Assembly with C

- Some things must be done in assembly
- Anything where you manipulate the stack

## Infrastructure in precode

- Extra `_asm.S` file for code units that need assembly

header	<code>scheduler.h</code>	<code>syscall.h</code>
C code	<code>scheduler.c</code>	<code>syscall.c</code>
assembly code	<code>scheduler_asm.S</code>	<code>syscall_asm.S</code>

- Include file: `asm_common.h.S`
  - Place to define asm macros
  - Recommend: asm `.macro` over C `#define`
- Export file: `asm-offsets.c`
  - Export C constants to asm
  - Especially struct field offsets  
e.g. `offsetof(struct pcb, next)`

## Task: Implement Locks: Blocking and Unblocking

- Locks in `sync.ch`
- We don't really *need* synchronization at this point
- Because threads can't be interrupted
- But these locks are here to test your scheduler
  - The lock-test and Monte-Carlo Pi threads use the lock API
  - If you mishandle the locks, they will not behave properly
- Can you block and unblock threads?
  - Not specific to locks
  - General service that the kernel should support

```

struct _lock {
    /* TODO: Design your lock struct.
     * You'll need at the very least a flag
     * and a wait queue. */
};

typedef struct _lock lock_t;

```

```

#define LOCK_INIT \
    { \
    }

void lock_acquire(lock_t *);
void lock_release(lock_t *);

```

### Task: Implement System Call Mechanism

- How does a user process *outside* of the kernel make a call *into* the kernel?
  - Kernel and process are separate executables
  - Not linked together
  - Do not know exact memory addresses of kernel functions
- Need some kind of agreed-upon protocol
  - Project 2 approach: Indirect function call via fixed address (0xf00)

#### USER SIDE

```

process    syslib    syslib
main() -> yield() -> invoke_syscall_entry_fixedpoint(YIELD) ->
                        -> pointer (*0xf00)(YIELD) ->

```

#### KERNEL SIDE

```

kernel asm          kernel C          kernel C
-> syscall_entry(YIELD) -> syscall_dispatch(YIELD) -> yield()

```

### Syscall vs Scheduler

- Both require low-level context switch
- Syscall:
  - same thread
  - switch from *user* context to *kernel* context
- Scheduler:
  - already in kernel context
  - switch from *one thread* to *another*
- Keeping this stuff straight can be a little mind-bending
  - Need to look behind abstractions like flow of execution and function calls
  - Think about raw program state: register values, stack pointer, instruction pointer
  - Low-level kernel code has to get very meta: manipulate state, shuffle stack pointers



### Task: Measure Context Switch Time

- You decide what to measure and how (methodology)
  - There is code in `util.ch` for measuring CPU time
- Do your measurements
- Get results
- Write it up in your report

### Extra Challenges

- We can't actually give extra credit on an assignment that doesn't count towards your grade
    - But trying these challenges will enhance your understanding
    - And make you stand out
    - Maybe get a TA job next year?
1. Implement something similar to the unix `time` command

```
$ time make clean all
real    0m3.454s
user    0m2.517s
sys     0m0.923s
```

    - Measure time in user space vs time in kernel space (“sys”)
    - Don't bother trying to calculate real-time
  2. More threads
    - Add a new in-kernel thread
    - Add a new user process

## Administrative Details

### Design Reviews

- Meet with your TA and present your design for the project
  - In colloquium period ~1 week after hand-out
  - TAs will schedule design review times
- This is mandatory. The design review is part of the assignment.
- This is not a formal presentation
  - But you should have some slides/visuals to show
  - Your task is to convince them that you understand the project
- Keep it at the design level
  - You don't need to go deep into implementation details
  - But we want to see that you have a ideas for implementation

### **Possible Topics for Design Review**

- What will go into your PCB? And why?
- How will you implement locks and process queues?
- What do you need to save on context switch?
  - Where will you save it?
- What is the difference between processes and threads?
  - How are the concepts muddled in our precode?
- How many stacks will you use per thread? And why?

### **Code**

- Code should be as readable as possible
  - Put thought into names and order
- Be sure to comment your code
  - Comments should explain the reasoning behind the code
  - This is especially important when dealing with entry points, synchronization, and blocking
  - Comments are part of the grading
- Be sure to test your code
  - Run your OS in the emulator
  - If writing code in `lib/`, you can try using the unit test framework

### **Report**

- Should be around 4 pages
- Give an overview of how you solved each task (or extra challenge)
- Describe how you tested your code
  - Point out any known bugs/issues
  - Describe how you would try to fix the bugs if you had more time
- Describe the methodology, results, and conclusions of your performance measurements

### **Hand In via Canvas**

- Put your report in your repository, under a `report/` dir
  - Report should be a PDF format
  - If you write in Word or other WYSIWYG word processor, export to PDF
- If you write in a document prep system like Markdown or LaTeX, you can include the report source if you like, but it's not required.

- Zip up your entire repository (code tree + report + .git/ dir)
- Submit via Canvas