UiT The Arctic University of Norway

Project 3: Preemptive Multitasking
Managing processes

INF-2201 Staff

UiT

Spring 2024

Project 3: Administrative Info (just like Project 2)

- ► Mandatory assignment
  - ► This is a mandatory assignment
  - ► It will be graded as pass/fail by TAs
  - ► You must pass to gain admission to the exam

- ► Groups of two
  - ► You must work in groups of two

- ► Design review
  - ► Meet with TA and present your design
  - ► Informal, but must have some kind of slides / presentation

- ► Hand-in
  - ► Hand-in via Canvas
  - ► Code: whole repository (working tree + .git/ dir)
  - ► Report: place in a report/ directory in your repo
  - ► Zip up and upload to Canvas

Project 3: Preemptive multitasking

- ▶ Project 2 was cooperative multitasking
  - ▶ Threads had to call `yield` or `exit` to give up control
- ▶ In Project 3: we will *preempt* the threads
  - ▶ Timer interrupt will fire every 10 ms
  - ▶ Give the kernel a chance to take over and switch threads
- ▶ You will:
  - ▶ Add timer interrupt handler code
  - ▶ Re-implement locks so that they work with preemption
  - ▶ Implement other sync abstractions: semaphores, condition variables, barriers
- ▶ Extra challenges:
  - ▶ Priority scheduling
  - ▶ Reimplement the dining philosophers on Linux with POSIX threads

Project 3: Environment / Precode

Same simple environment

- ▶ Protected mode, but no active protection
    - ▶ CPU in 32-bit Protected Mode
    - ▶ No protection active: all runs at kernel level (Ring 0)
    - ▶ Flat 32-bit address space
    - ▶ Kernel and processes share one address space
- ▶ No `malloc` / `free`
    - ▶ Statically allocate any globals / arrays you need
    - ▶ Allocate single structs from arrays

What's new?

- ▶ Precode includes our Project 2 solution, plus…
- ▶ Interrupts!
    - ▶ Interrupts are turned on
    - ▶ Syscall entry via interrupt
- ▶ Real time emulation
    - ▶ Bochs time syncs with real time
    - ▶ OS detects CPU speed with timer

Interrupts

Triggers

- ▶ Signal from hardware: timer, keyboard, etc.
- ▶ CPU exception: divide by zero, invalid memory access, etc.
- ▶ Software trigger: INT instruction

Operation

- ▶ CPU saves state to stack and jumps to pre-set address
- ▶ Like a function call that can happen at any time
- ▶ *Interrupt Handler* aka *Interrupt Service Routine (ISR)*

Saved state

- ▶ Not just EIP: EFLAGS, CS, EIP
- ▶ EFLAGS Interrupt Flag (IF) is cleared, to disable interrupts during handler
- ▶ Special return instruction: `iret`

Interrupt Hardware Support
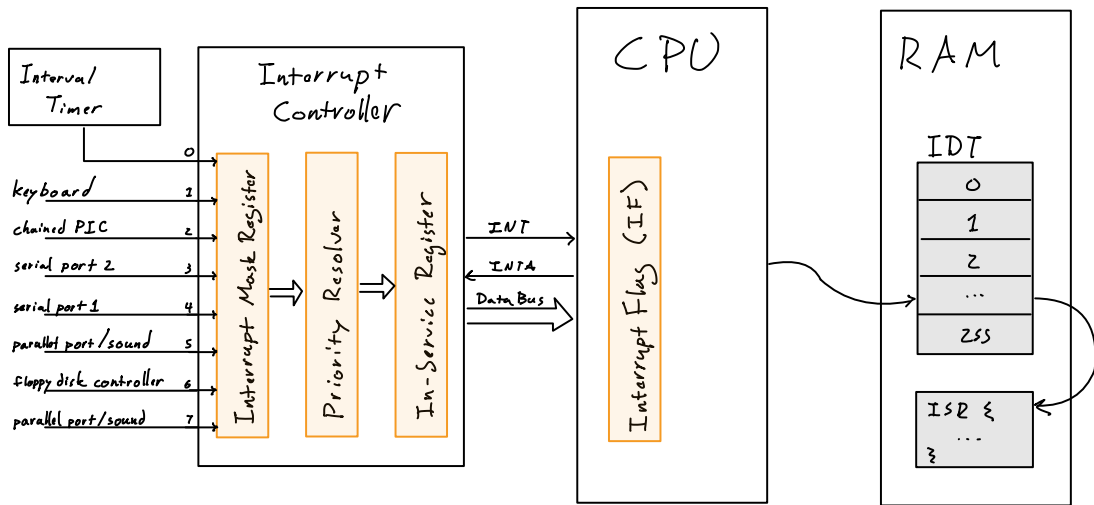
Interrupts, from hardware to CPU

Interrupt Hardware Support

- ▶ Devices make *Interrupt Requests (IRQs)*

    - ▶ 0 = timer
    - ▶ 1 = keyboard
    - ▶ …

- ▶ to *Programmable Interrupt Controller (PIC)*

    - ▶ PIC multiplexes IRQs and sends to CPU
    - ▶ Maps IRQ (0–7) to *interrupt vector* (0–255)
    - ▶ Typically IRQ 0 maps to interrupt 32

- ▶ to CPU

    - ▶ Gets signal and vector number
    - ▶ Looks in *Interrupt Descriptor Table (IDT)*
    - ▶ Descriptor points to ISR
    - ▶ CPU Executes ISR
    - ▶ ISR `iret` returns to interrupted execution
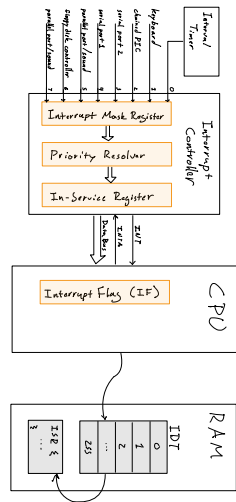
Interrupts, from hardware to CPU

# Three Chances to Block an Interrupt

1. PIC: Interrupt Mask Register

   ▶ CPU can command PIC to ignore specific IRQs

2. PIC: In-Service Register

   ▶ When PIC sends interrupt to CPU,
     it marks the IRQ as *in service*.
   ▶ Will not send the same interrupt again,
     until the CPU sends an *End Of Interrupt (EOI)*.

3. CPU: Interrupt Flag (IF)

   ▶ In EFLAGS register
   ▶ Clear to ignore all external interrupts
   ▶ CPU clears flag at interrupt start
   ▶ Flag restored when EFLAGS is restored on `iret`
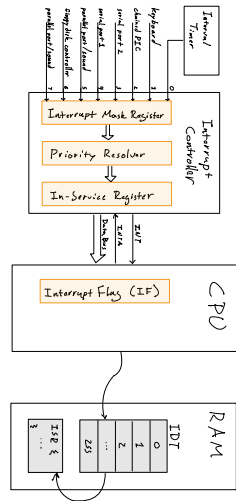
Interrupts, from hardware to CPU

Interrupt Sequence

1. IRQ comes in to PIC

   1.1 Check interrupt mask
   1.2 Check in-service
   1.3 Mark IRQ as in service
       (disable this IRQ)
   1.4 Send interrupt vector to CPU

2. CPU gets interrupt

   2.1 Check Interrupt Flag
   2.2 Look up interrupt in IDT
   2.3 Push EFLAGS, CS, IP
   2.4 Disable Interrupt Flag
       (disable interrupts)
   2.5 Jump to beginning of ISR

3. Interrupt Service Routine

   3.1 Save additional context
   3.2 Actually handle interrupt
   3.3 Restore additional context
   3.4 Send EOI (re-enable this IRQ)
   3.5 `iret`

4. CPU returns

   4.1 Pop and restore EIP, CS, EFLAGS
       (re-enable interrupts)
   4.2 Resume execution

Interrupt on a Timer for Preemptive Scheduling

- ▶ *Programmable Interval Timer (PIT)* chip
    - ▶ Base clock ticks at 1193180 Hz (1.19 MHz)
    - ▶ Output channel can fire at any fraction of that
    - ▶ Output channel 0 is connected to PIC IRQ 0

- ▶ Precode:
    - ▶ Sets up PIT: Fire IRQ 0 every 10 ms
    - ▶ Sets up PIC: Map IRQ 0 to interrupt 32, mask others
    - ▶ Sets up IDT: Handler for interrupt 32 → timer_isr_entry

- ▶ You: Implement timer_isr_entry function (in interrupt_asm.S)
    - ▶ Save/restore necessary context
    - ▶ Switch to kernel stack if using user stack
    - ▶ Call into scheduler to switch to next process
    - ▶ Switch back to user stack if necessary
    - ▶ Restore necessary context
    - ▶ Send EOI

Tasks Overview

1. Implement preemptive scheduling

2. Implement synchronization abstractions

   2.1 Re-implement locks to work with preemption
   2.2 Implement semaphores
   2.3 Implement condition variables
   2.4 Implement barriers

3. Ensure that preemption does not break existing `yield` and `exit` calls

4. Fix the dining philosophers implementation

Task: Implement preemptive scheduling

▶ Need to implement `timer_isr_entry` function (in interrupt_asm.S)

  ▶ Save/restore necessary context
  ▶ Switch to kernel stack if using user stack
  ▶ Call into scheduler to switch to next process
  ▶ Switch back to user stack if necessary
  ▶ Restore necessary context
  ▶ Send EOI

▶ May need to add/change some code in scheduler as well

Task: Re-Implement Locks

1. Re-implement locks to work with preemption

   ▶ Remember: timer interrupt can happen between any two instructions
   ▶ Need some kind of atomicity

▶ Provided helper functions for "no-interrupt" sections

   ▶ Call `nointerrupt_enter` to disable interrupts
   ▶ Call `nointerrupt_leave` to reenable interrupts
   ▶ Can be nested, will only re-enable on last leave

▶ Precode includes skeletons for multiple implementations:

   ▶ Project 2 version: `_coop` suffix (don't bother)
   ▶ Project 3 version: `_nointerrupt` suffix (impement these)

Task: Implement Other Sync Abstractions

2. Implement semaphores
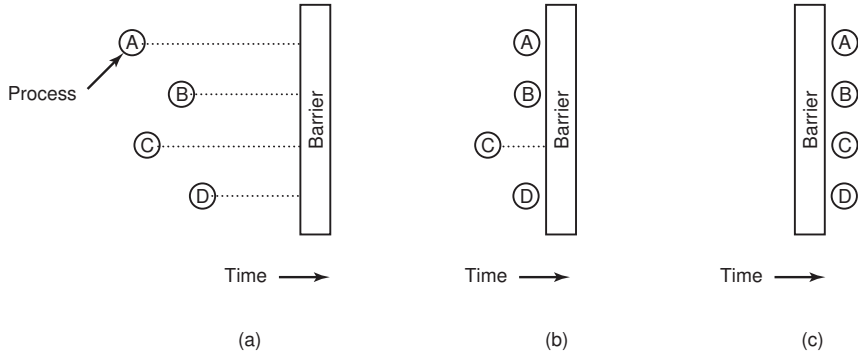
   ▶ Test code: dining philosophers (philosophers.c)

3. Implement condition variables (monitors)

   ▶ Test code: lock test threads (th2.c)
     upgraded to also use condvars

4. Implement barriers

   ▶ Test code: barrier test threads (barrier_test.c)

Barriers

Figure 2-37. Use of a barrier. (a) Processes approaching a barrier. (b) All processes but one blocked at the barrier. (c) When the last process arrives at the barrier, all of them are let through.

Barriers (Tannenbaum, *Modern Operating Systems*, 4th ed)

Task: Ensure that preemption does not break `yield` and `exit`

▶ Precode includes syscall implementations
▶ New implementation uses interrupt (INT 48) instead of fixed address
▶ Precode includes scheduler implementation
▶ You may need to make some changes to ensure that `yield` and `exit` still work with preemption active
▶ Or they're fine and you just need to make sure you don't break them

Task: Fix the dining philosophers implementation

- ▶ Precode includes an implementation of the Dining Philosophers problem

- ▶ Display on screen and also keyboard LEDs: num lock, caps lock, scroll lock

- ▶ Our solution is not fair
  - ▶ The middle philosohper, Caps Lock, as a slight advantage
  - ▶ Over time, he gets to eat more than the others
  - ▶ Why is this?

- ▶ Fix the philosophers code so that it is fair
  - ▶ Be careful not to introduce deadlocks or starvation
  - ▶ Explain your solution in your report
  - ▶ Is it fair? How do you know
  - ▶ Are there deadlocks or starvation? How do you know?

Extra Challenges

1. Implement priority scheduling

   ▶ The plane process (process1.c) uses the new `setpriority` syscall
   ▶ Implement priority scheduling: should see the plane's speed change

2. Re-implement dining philosophers on Linux using POSIX threads (pthreads)

   ▶ There is no precode for this
   ▶ Set up a simple C project and write your own dining philosphers simulation

Where to start

- ▶ Implement the preemptive scheduling first
- ▶ Then the lock implementation
- ▶ Then semaphores, then condition variables, then barriers
- ▶ Remember: all sync abstractions should work with preemption
- ▶ Finally, fix the philosophers

Procedure is the same as Project 2

▶ Design reviews
▶ Code
▶ Report
▶ Hand in via canvas: zip up entire repository plus report

Design Reviews

- ▶ Design reviews start next week
    - ▶ For some, this means Monday (!!)

- ▶ Remember: informal presentation
    - ▶ Don't need code
    - ▶ Need some kind of visual
    - ▶ Need to show that you understand the theory and that you've got a plan

- ▶ Keep it at the design level
    - ▶ You don't need to go deep into implementation details
    - ▶ But we want to see that you have a ideas for implementation

- ▶ Notes to individual groups:
    - ▶ Group 1 with Øyvind: Most DRs will be online (Discord)
    - ▶ Group 2 with Ilya: Friday will be a presentation about how to write a good report

Report

- ▶ Should be around 4 pages
- ▶ Give an overview of how you solved each task (or extra challenge)
- ▶ Describe how you tested your code
    - ▶ Point out any known bugs/issues
    - ▶ Describe how you would try to fix the bugs if you had more time
- ▶ Describe the methodology, results, and conclusions of your performance measurements
- ▶ We are working on a "How to Write a Report" guide and report template
    - ▶ Expect those late this week or early next week

Hand In via Canvas

- ▶ Put your report in your repository, under a report/ dir
    - ▶ Report shoud be a PDF format
    - ▶ If you write in Word or other WYSIWYG word processor, export to PDF
- ▶ Zip up your entire repository (code tree + report + .git/ dir)
- ▶ Submit via Canvas