



UiT The Arctic University of Norway

# Project 3: Preemptive Multitasking

Managing processes

INF-2201 Staff

UiT

Spring 2024

## Project 3

Project Overview

Interrupts

Project Tasks

Administrative Details

# Project 3: Administrative Info (just like Project 2)

- ▶ Mandatory assignment
  - ▶ This is a mandatory assignment
  - ▶ It will be graded as pass/fail by TAs
  - ▶ You must pass to gain admission to the exam
- ▶ Groups of two
  - ▶ You must work in groups of two
- ▶ Design review
  - ▶ Meet with TA and present your design
  - ▶ Informal, but must have some kind of slides / presentation
- ▶ Hand-in
  - ▶ Hand-in via Canvas
  - ▶ Code: whole repository (working tree + `.git/` dir)
  - ▶ Report: place in a `report/` directory in your repo
  - ▶ Zip up and upload to Canvas

# Project 3: Preemptive multitasking

- ▶ Project 2 was cooperative multitasking
  - ▶ Threads had to call `yield` or `exit` to give up control
- ▶ In Project 3: we will *preempt* the threads
  - ▶ Timer interrupt will fire every 10 ms
  - ▶ Give the kernel a chance to take over and switch threads
- ▶ You will:
  - ▶ Add timer interrupt handler code
  - ▶ Re-implement locks so that they work with preemption
  - ▶ Implement other sync abstractions: semaphores, condition variables, barriers
- ▶ Extra challenges:
  - ▶ Priority scheduling
  - ▶ Reimplement the dining philosophers on Linux with POSIX threads

# Project 3: Environment / Precode

## Same simple environment

- ▶ Protected mode, but no active protection
  - ▶ CPU in 32-bit Protected Mode
  - ▶ No protection active: all runs at kernel level (Ring 0)
  - ▶ Flat 32-bit address space
  - ▶ Kernel and processes share one address space
- ▶ No `malloc` / `free`
  - ▶ Statically allocate any globals / arrays you need
  - ▶ Allocate single structs from arrays

## What's new?

- ▶ Precode includes our Project 2 solution, plus...
- ▶ Interrupts!
  - ▶ Interrupts are turned on
  - ▶ Syscall entry via interrupt
- ▶ Real time emulation
  - ▶ Bochs time syncs with real time
  - ▶ OS detects CPU speed with timer

## Project 3

Project Overview

**Interrupts**

Project Tasks

Administrative Details

# Interrupts

## Triggers

- ▶ Signal from hardware: timer, keyboard, etc.
- ▶ CPU exception: divide by zero, invalid memory access, etc.
- ▶ Software trigger: INT instruction

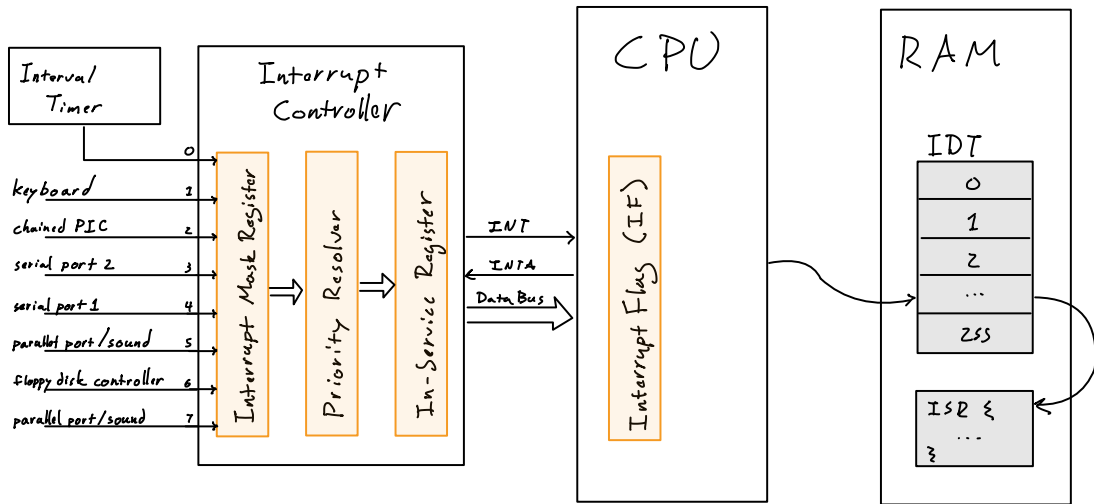
## Operation

- ▶ CPU saves state to stack and jumps to pre-set address
- ▶ Like a function call that can happen at any time
- ▶ *Interrupt Handler* aka *Interrupt Service Routine (ISR)*

## Saved state

- ▶ Not just EIP: EFLAGS, CS, EIP
- ▶ EFLAGS Interrupt Flag (IF) is cleared, to disable interrupts during handler
- ▶ Special return instruction: `iret`

# Interrupt Hardware Support

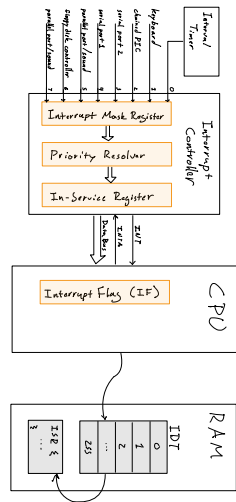


Interrupts, from hardware to CPU



# Interrupt Hardware Support

- ▶ Devices make *Interrupt Requests (IRQs)*
  - ▶ 0 = timer
  - ▶ 1 = keyboard
  - ▶ ...
- ▶ to *Programmable Interrupt Controller (PIC)*
  - ▶ PIC multiplexes IRQs and sends to CPU
  - ▶ Maps IRQ (0-7) to *interrupt vector* (0-255)
  - ▶ Typically IRQ 0 maps to interrupt 32
- ▶ to CPU
  - ▶ Gets signal and vector number
  - ▶ Looks in *Interrupt Descriptor Table (IDT)*
  - ▶ Descriptor points to ISR
  - ▶ CPU Executes ISR
  - ▶ ISR `iret` returns to interrupted execution



Interrupts, from hardware to CPU

# Three Chances to Block an Interrupt

## 1. PIC: Interrupt Mask Register


- ▶ CPU can command PIC to ignore specific IRQs

## 2. PIC: In-Service Register

- ▶ When PIC sends interrupt to CPU, it marks the IRQ as *in service*.
- ▶ Will not send the same interrupt again, until the CPU sends an *End Of Interrupt (EOI)*.

## 3. CPU: Interrupt Flag (IF)

- ▶ In EFLAGS register
- ▶ Clear to ignore all external interrupts
- ▶ CPU clears flag at interrupt start
- ▶ Flag restored when EFLAGS is restored on `iret`



Interrupts, from hardware to CPU

# Interrupt Sequence

## 1. IRQ comes in to PIC

- 1.1 Check interrupt mask
- 1.2 Check in-service
- 1.3 Mark IRQ as in service  
(disable this IRQ)
- 1.4 Send interrupt vector to CPU

## 2. CPU gets interrupt

- 2.1 Check Interrupt Flag
- 2.2 Look up interrupt in IDT
- 2.3 Push EFLAGS, CS, IP
- 2.4 Disable Interrupt Flag  
(disable interrupts)
- 2.5 Jump to beginning of ISR

## 3. Interrupt Service Routine

- 3.1 Save additional context
- 3.2 Actually handle interrupt
- 3.3 Restore additional context
- 3.4 Send EOI (re-enable this IRQ)
- 3.5 `iret`

## 4. CPU returns

- 4.1 Pop and restore EIP, CS, EFLAGS  
(re-enable interrupts)
- 4.2 Resume execution

# Interrupt on a Timer for Preemptive Scheduling

- ▶ *Programmable Interval Timer (PIT)* chip
  - ▶ Base clock ticks at 1193180 Hz (1.19 MHz)
  - ▶ Output channel can fire at any fraction of that
  - ▶ Output channel 0 is connected to PIC IRQ 0
- ▶ Precode:
  - ▶ Sets up PIT: Fire IRQ 0 every 10 ms
  - ▶ Sets up PIC: Map IRQ 0 to interrupt 32, mask others
  - ▶ Sets up IDT: Handler for interrupt 32 → `timer_isr_entry`
- ▶ You: Implement `timer_isr_entry` function (in `interrupt_asm.S`)
  - ▶ Save/restore necessary context
  - ▶ Switch to kernel stack if using user stack
  - ▶ Call into scheduler to switch to next process
  - ▶ Switch back to user stack if necessary
  - ▶ Restore necessary context
  - ▶ Send EOI

## Project 3

Project Overview

Interrupts

**Project Tasks**

Administrative Details

# Tasks Overview

1. Implement preemptive scheduling
2. Implement synchronization abstractions
  - 2.1 Re-implement locks to work with preemption
  - 2.2 Implement semaphores
  - 2.3 Implement condition variables
  - 2.4 Implement barriers
3. Ensure that preemption does not break existing `yield` and `exit` calls
4. Fix the dining philosophers implementation

# Task: Implement preemptive scheduling

- ▶ Need to implement `timer_isr_entry` function (in `interrupt_asm.S`)
  - ▶ Save/restore necessary context
  - ▶ Switch to kernel stack if using user stack
  - ▶ Call into scheduler to switch to next process
  - ▶ Switch back to user stack if necessary
  - ▶ Restore necessary context
  - ▶ Send EOI
- ▶ May need to add/change some code in scheduler as well

# Task: Re-Implement Locks

## 1. Re-implement locks to work with preemption

- ▶ Remember: timer interrupt can happen between any two instructions
- ▶ Need some kind of atomicity
- ▶ Provided helper functions for “no-interrupt” sections
  - ▶ Call `nointerrupt_enter` to disable interrupts
  - ▶ Call `nointerrupt_leave` to reenable interrupts
  - ▶ Can be nested, will only re-enable on last leave
- ▶ Precode includes skeletons for multiple implementations:
  - ▶ Project 2 version: `_coop` suffix (don't bother)
  - ▶ Project 3 version: `_nointerrupt` suffix (implement these)



# Task: Implement Other Sync Abstractions

## 2. Implement semaphores

- ▶ Test code: dining philosophers (philosophers.c)


## 3. Implement condition variables (monitors)

- ▶ Test code: lock test threads (th2.c)  
upgraded to also use condvars

## 4. Implement barriers

- ▶ Test code: barrier test threads (barrier\_test.c)

# Barriers



**Figure 2-37.** Use of a barrier. (a) Processes approaching a barrier. (b) All processes but one blocked at the barrier. (c) When the last process arrives at the barrier, all of them are let through.

Barriers (Tannenbaum, *Modern Operating Systems*, 4th ed)

# Task: Ensure that preemption does not break `yield` and `exit`

- ▶ Precode includes syscall implementations
- ▶ New implementation uses interrupt (INT 48) instead of fixed address
- ▶ Precode includes scheduler implementation
- ▶ You may need to make some changes to ensure that `yield` and `exit` still work with preemption active
- ▶ Or they're fine and you just need to make sure you don't break them

# Task: Improve the Dining Philosophers implementation

- ▶ Precode includes an implementation of the Dining Philosophers problem
- ▶ Display on screen and also keyboard LEDs: num lock, caps lock, scroll lock
- ▶ Our solution is not fair
  - ▶ The middle philosopher, Caps Lock, as a slight advantage
  - ▶ Over time, he gets to eat more than the others
  - ▶ Why is this?
- ▶ Analyze this in your report. Answer these questions:
  - ▶ Can this solution deadlock?
  - ▶ What makes this solution unfair?
  - ▶ How does this relate to the concept of *starvation*?
  - ▶ How would you go about making it fair?
- ▶ Update the existing solution: try to make it more fair

# Extra Challenges

## 1. Implement priority scheduling

- ▶ The plane process (process1.c) uses the new `setpriority` syscall
- ▶ Implement priority scheduling: should see the plane's speed change

## 2. Re-implement dining philosophers on Linux using POSIX threads (pthreads)

- ▶ There is no precode for this
- ▶ Set up a simple C project and write your own dining philosophers simulation

# Where to start

- ▶ Implement the preemptive scheduling first
- ▶ Then the lock implementation
- ▶ Then semaphores, then condition variables, then barriers
- ▶ Remember: all sync abstractions should work with preemption
- ▶ Finally, fix the philosophers

## Project 3

Project Overview

Interrupts

Project Tasks

Administrative Details

# Procedure is the same as Project 2

- ▶ Design reviews
- ▶ Code
- ▶ Report
- ▶ Hand in via canvas: zip up entire repository plus report



# Design Reviews

- ▶ Design reviews start next week
  - ▶ For some, this means Monday (!!)
- ▶ Remember: informal presentation
  - ▶ Don't need code
  - ▶ Need some kind of visual
  - ▶ Need to show that you understand the theory and that you've got a plan
- ▶ Keep it at the design level
  - ▶ You don't need to go deep into implementation details
  - ▶ But we want to see that you have a ideas for implementation
- ▶ Notes to individual groups:
  - ▶ Group 1 with Øyvind: Most DRs will be online (Discord)
  - ▶ Group 2 with Ilya: Friday will be a presentation about how to write a good report

# Report

- ▶ Should be around 4 pages
- ▶ Give an overview of how you solved each task (or extra challenge)
- ▶ Describe how you tested your code
  - ▶ Point out any known bugs/issues
  - ▶ Describe how you would try to fix the bugs if you had more time
- ▶ Describe the methodology, results, and conclusions of your performance measurements
- ▶ We are working on a “How to Write a Report” guide and report template
  - ▶ Expect those late this week or early next week

# Hand In via Canvas

- ▶ Put your report in your repository, under a `report/` dir
  - ▶ Report should be a PDF format
  - ▶ If you write in Word or other WYSIWYG word processor, export to PDF
- ▶ Zip up your entire repository (code tree + report + `.git/` dir)
- ▶ Submit via Canvas