

Project 4: Inter-Process Communication

INF-2201 Staff

Spring 2024

Contents

Project 4	1
Project Overview	1
Project 4: Administrative Info (mostly like the others)	1
Project 4: Inter-Process Communication (IPC)	2
Project 4: Environment / Precode	2
Project Tasks	3
Tasks Overview	3
Task: Implement IPC Message Passing	3
Mailbox API	3
Task: Complete Keyboard Driver	4
Keyboard Interface with OS: putchar	4
Keyboard: Subtle Points	5
Keyboard: Deadlock Danger	5
Task: Reduce Interrupt Disable Time	5
Rewriting Sync Abstractions	6
Using Spinlocks: An Example	6
Task: Load Processes Dynamically	7
Loading Need 1: Separate Address Spaces	7
Loading Need 2: Memory Manager	7
Loading Need 3: Disk Format that Lists Processes	8
Implementing Loading	8
Implementation Help: "Given" Files	9
Extra Challenges	9
Administrative Details	10
Procedure is (Mostly) the Same as Other Projects	10
Report	10
Hand In via Canvas	10

Project 4

Project Overview

Project 4: Administrative Info (mostly like the others)

- Mandatory assignment

- Groups of two
- Design review
- Hand-in
 - Hand-in via Canvas
 - Code: whole repository (working tree + `.git/` dir)
 - Report: place in a `report/` directory in your repo
 - Zip up and upload to Canvas
 - **New!** Also upload the report PDF separately

Project 4: Inter-Process Communication (IPC)

- In Project 3:
 - We had separation between processes and kernel
 - But we didn't enforce it
- In Project 4:
 - We will enforce the separation
 - But we will build an Inter-Process Communication system
- You will:
 - Implement IPC message passing
 - Complete a keyboard driver
 - Reimplement sync abstractions to reduce interrupt disable time
 - Load processes dynamically

Project 4: Environment / Precode

Precode includes our solution to Project 3, plus many additions to support Project 4.

Project 3	Project 4
All kernel priv (0)	Priv enforced: Kernel (0) vs Process (3)
One address space	Each process has own address space
Atomic via disabling interrupts	Atomic via C <code><stdatomic.h></code>
No interaction	Keyboard input + interactive shell
No disk I/O after boot	USB storage drivers
Flat disk image	Image has simple directory of processes

Be warned: Later projects have received less attention

- You will encounter more old, un-refactored code
- There will be bugs

Project Tasks

Tasks Overview

You will:

- Implement IPC message passing
- Complete keyboard driver
- Reduce interrupt disable time (by reimplementing sync abstractions)
- Load processes dynamically

Task: Implement IPC Message Passing




Figure 1: A ring-buffer as a mailbox

Ring buffer

- Fixed-size: *bounded buffer*
- Variable-sized messages
- First In, First Out (FIFO)

Producer-consumer problem

- Multiple producers
- Multiple consumers
- Blocking operations

Implementation

- `mbox.[ch]`
- Struct already defined
- Locks + condition variables
- You will write the API functions

Mailbox API

- API functions

```
void mbox_init(void);  
int mbox_open(int key);
```

```

int mbox_close(int q);
int mbox_stat(int q, int *count, int *space);
int mbox_recv(int q, msg_t *m);
int mbox_send(int q, msg_t *m);

```

- Mailboxes identified by integer id (like file descriptors)
- No other addressing
- Any process can put a message in a box
- Any process can remove a message from a box

Task: Complete Keyboard Driver

Chain

IRQ 1	->	keyboard_isr_entry	->	keyboard_interrupt	->	key handlers	->	putchar
(hw)		(interrupt_asm.S)		(keyboard.c)		(keyboard.c)		(keyboard.c)
-----							-----	
Provided by precode							Your task	

Hardware

- Hardware IRQ 1
- Read *scan code* from port 0x60

Partial driver

- keyboard_interrupt: look up scan code in table
- table has pointers to handler fns
- handler fns:
 - test for CTRL / SHIFT / ALT
 - map scan code to ASCII
 - call putchar

Keyboard Interface with OS: putchar

Precode

- Maps keyboard scan code to ASCII
- NB: US keyboard layout

putchar

- Put character into mailbox
- How exactly? Up to you

Classic producer/consumer problem

- Single producer. Or is it multiple?
- One source of input: keyboard
- But interrupt occurs in context of different threads/processes

- Multiple consumers. Or is it single?
- Any process can pull ASCII messages off of the mailbox
- But in practice it will usually just be the shell

Keyboard: Subtle Points

- Producer must not block
 - This is important for a keyboard handler
 - Why is this?
 - Solution: drop key if buffer is full
 - Check before send: `mbox_stat` before `mbox_send`
 - Need check + send to be atomic. Why?
- What if `getchar` is interrupted by a keyboard interrupt?
 - Why is this a problem?
 - Solution: disable interrupt. But where?
- ↑ Design review questions ↑

Keyboard: Deadlock Danger

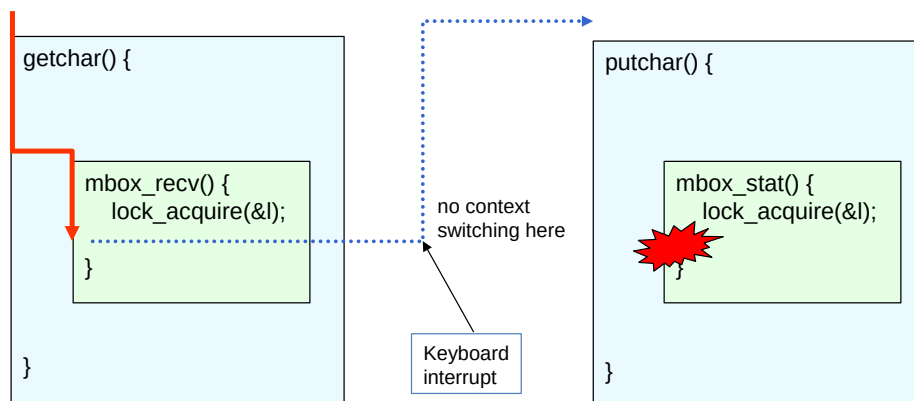


Figure 2: `getchar` interrupted by keyboard interrupt

Task: Reduce Interrupt Disable Time

- Interrupts should be disabled as little as possible
 - Potential to lose hardware events
 - Interrupt controller will wait for CPU, but...
 - What if another key comes in before you handle the first?
 - What if another timer event fires before you handle the last?
- Where can we reduce?
 - Scheduler? Too difficult to rewrite this time
 - Sync abstractions? ...

Rewriting Sync Abstractions

- In Project 3, we achieved atomicity by disabling interrupts
- Alternative: atomic test-and-set
 - C standard <stdatomic.h> provides `atomic_flag` type
 - `atomic_flag_test_and_set()`
 - `atomic_flag_clear()`
 - Compile atomic CPU instructions

- Use atomic-test-and-set to build a spinlock

```
spinlock_acquire(atomic_flag *flag) {  
    while(atomic_test_and_set(flag)) yield();  
}
```

```
spinlock_release(atomic_flag *flag) {  
    atomic_flag_clear(flag);  
}
```

- Use spinlocks to implement locks, condvars, etc.

Using Spinlocks: An Example

Via interrupt disable

```
static void lock_acquire(lock_t *l)  
{  
    nointerrupt_enter();  
    while (l->locked) {  
        block(&l->wait_queue);  
    }  
    l->locked = true;  
    nointerrupt_leave();  
}
```

Via spinlock

```
static void lock_acquire(lock_t *l)  
{  
    spinlock_acquire(&l->spinlock);  
    while (l->locked) {  
        block(&l->wait_queue, &l->spinlock);  
    }  
    l->locked = true;  
    spinlock_release(&l->spinlock);  
}
```

- `block` takes lock param
 - Releases before blocking
 - Re-acquires before returning
- This is no longer in the precode

- Should it be? Possible regression?

Task: Load Processes Dynamically

- In Project 3, all programs loaded at startup
- In Project 4, we will load them dynamically
- For dynamic loading, we need:
 1. Separate address space for each process
 2. Memory manager
 3. Disk format that lists processes
 - These are provided in precode

You will implement...

`readdir()`

1. Load process-directory sector
2. Return data to shell for parsing

`loadproc()`

1. Allocate physical memory
2. Read code/data from disk
3. Pass mem info to `create_process()`

Loading Need 1: Separate Address Spaces

- Privilege enforcement via Global Descriptor Table
 - Descriptors for kernel-level code/data (ring 0)
 - Descriptors for user-level code/data (ring 3)
 - Precode `pcb` struct now has CS/DS fields
 - In-kernel threads: CS/DS point to kernel-level descriptors
 - User processes: CS/DS point to user-level descriptors
- Separate address spaces for user processes
 - Process loaded to virtual address 0x100 0000
 - GDT Task State Segment to help with task switching
 - Precode will handle the details
 - You just need to implement physical page allocation
 - In Project 5 you will get into the details

Loading Need 2: Memory Manager

- Pool of physical pages
 - 4 KiB per page
 - 1 MiB pool

- You implement: `alloc_memory()`
 - Get pages from the pool
- Extra challenge: implement `free_memory()`

Loading Need 3: Disk Format that Lists Processes




Figure 3: Updated disk format with process directory

- Add a process directory after the kernel
- Very simple process directory format:
 - List of (location, size) pairs, both `int`
 - No names. Processes identified by index
 - Process 0: shell
 - Process 1: process1 (plane)
 - ...
- Kernel will load processes dynamically
 - Initiated by shell
 - `load 1: load process1 (plane)`
 - `load 2: load process2 (math)`
 - ...
- In diagram: "Program disk"?
 - Old feature to hot-swap floppies
 - No longer working

Implementing Loading

USB storage support (precode)

- Precode provides a basic USB storage driver (`kernel/usb/**`)
- Key header for you: `scsi.h`

- `scsi_read()`: reads one 512-byte sector
- `scsi_write()`: writes one 512-byte sector
- For now, only reading

Loading a program: `loadproc()` (you)

- Read process directory: `readdir()` (you)
- Allocate pages for code + data: `alloc_memory()` (you)
- Also allocate pages for stack
- Read process sectors: `scsi_read()` (precode)
- Initialize a `pcb` struct: `create_process()` (precode)

Implementation Help: “Given” Files

- This project has many pieces that lean on each other
 - Hard to develop and test one without the others
- “Given” files: Compiled binaries of working implementations:
 - `mbox.given.o`: IPC mailboxes
 - `keyboard.given.o`: keyboard driver
 - `sync.given.o`: sync abstractions
 - `pcb.given.o`: dynamic loading
- Lines in `Makefile.common`: Uncomment to use given files

```

KERNEL_OBJS := $(KERNEL_OBJS:kernel/mbox.o=kernel/mbox.given.o)
KERNEL_OBJS := $(KERNEL_OBJS:kernel/keyboard.o=kernel/keyboard.given.o)
KERNEL_OBJS := $(KERNEL_OBJS:kernel/sync.o=kernel/sync.given.o)
KERNEL_OBJS := $(KERNEL_OBJS:kernel/pcb.o=kernel/pcb.given.o)

```

- Careful with headers: Some changes may break given files

Extra Challenges

- Free memory when a process exits
- Use a more interesting allocation algorithm
 - Make it more like a proper `malloc`
- Implement new commands in the shell:
 - `ps`: list running processes
 - `kill`: stop a running process
 - Suspend / resume a process
- To really go above and beyond, get creative:
 - You have input now. You could write a simple game
 - Try switching to VGA graphics mode
 - Write a better shell

Administrative Details

Procedure is (Mostly) the Same as Other Projects

- Design reviews start on Monday (!!)
- Code
- Report
- Hand in via Canvas: zip up entire repository plus report
- **New!** Also upload the report PDF separately

Report

- Structure: scientific paper
- Length: around 4 pages
- Citations: required
 - You must cite sources you use
 - At the very least, you should have the textbook as a source
- File format: PDF
- AI tools: discouraged but not banned. Use must be declared
- **New!** See Howto doc and template in repository
 - `doc/how-to-write-a-report.md`: gives more guidelines and advice
 - `report/latex-src/template.tex`: LaTeX tutorial / template
 - If you use the template, be sure to remove all existing content

Hand In via Canvas

- Put your report in your repository, under a `report/` dir
 - Report should be a PDF format
 - If you write in Word or other WYSIWYG word processor, export to PDF
- If you write in a document prep system like Markdown or LaTeX, you can include the report source if you like, but it's not required.
- Zip up your entire repository (code tree + report + `.git/` dir)
- Submit via Canvas
 - Upload zip
 - **New!** Also upload report PDF
 - Having both makes it easier for us to grade