

Project 5: Virtual Memory

INF-2201 Staff

Spring 2024

Contents

Project 5	1
Project Overview	1
Project 5: Administrative Info (mostly like the others)	1
Project 5: Virtual Memory	2
Virtual Memory	2
Virtual Memory Overview	2
Multiple Address Spaces	3
x86-32 Page Mapping	3
x86-32 Page Mapping: More Info	3
x86-32 Page Mapping: MORE Info	4
Project Tasks	5
Tasks/Precode Overview	5
Task: Set Up New Process's Address Space	5
Task: Handle Page Faults (Swap)	5
Subtask: Physical Page Allocation	6
Extra Challenges	6
Hints	7
Administrative Details	7
Procedure is (Mostly) the Same as Other Projects	7
Report	7
Hand In via Canvas	7

Project 5

Project Overview

Project 5: Administrative Info (mostly like the others)

- Mandatory assignment
- Groups of two
- Design review
- Hand-in
 - Hand-in via Canvas

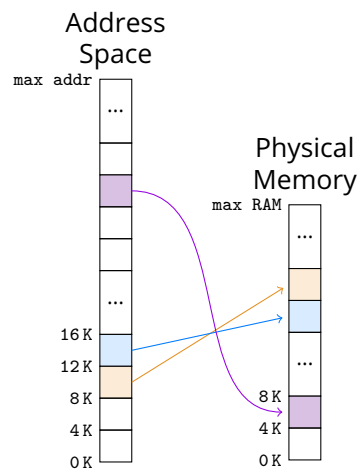


Figure 1: Mapping from virtual address space to physical memory

- Code: whole repository (working tree + .git/ dir)
- Report: place in a `report/` directory in your repo
- Zip up and upload to Canvas
- **New!-ish** Also upload the report PDF separately

Project 5: Virtual Memory

- In Project 4:
 - We started enforcing user/kernel separation
 - Processes had separate address spaces
 - Virtual memory turned on, but very basic
- In Project 5:
 - We will improve the virtual memory
 - Swap memory pages on demand
- You will:
 - Manage virtual memory with dynamic swapping
 - Manage page tables
 - Implement a page fault handler
 - Implement page allocation scheme
 - Swap to/from disk

Virtual Memory

Virtual Memory Overview

- Virtual memory
 - Add indirection!
 - Map virtual addresses → physical memory




Figure 2: Processes with separate address spaces

- Divide memory into *pages*
 - Common size: 4 KiB
 - 32-bit address space:
 $4 \text{ GiB space} = 4 \text{ KiB/page} \times 2^{20} \text{ pages/space}$
- Benefits
 - Use whole address space!
 - Swap pages to disk as needed!

Multiple Address Spaces

- Each process can have its own address space
- Great way to enforce process separation
 - Can't tamper with what you can't address
 - Can have per-page permissions
- But can map some pages into multiple spaces
 - Common to share kernel pages

x86-32 Page Mapping

x86-32 Page Mapping: More Info

- To map a 32-bit address space:
 - $4 \text{ GiB space} = 4 \text{ KiB/page} \times 2^{20} \text{ pages/space}$
 - 20 bits: can select page in space
 - 12 bits: can select byte in page
- Two-level hierarchy
 - Level 2: Page Directory
 - Level 1: Page Table

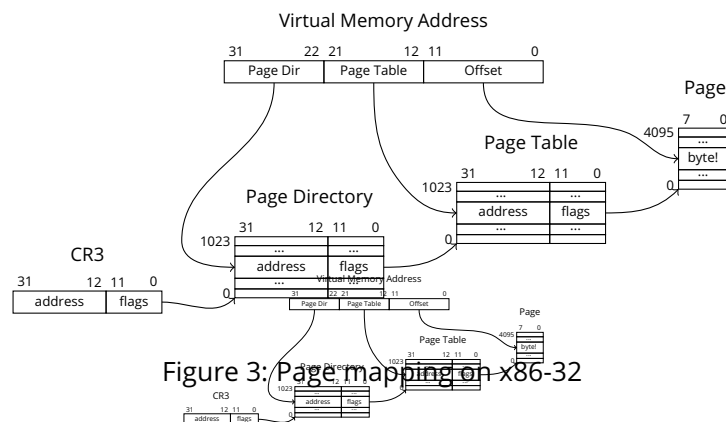



Figure 4: Page mapping on x86-32

- Level 0: Physical Page
- Each step in hierarchy is also a page
 - 4 KiB page = 32 bits/entry \times 1024 entries/page
 - 20 bits: phys addr of next page in hierarchy
 - 12 bits: flags
- CR3 \rightarrow 1 page dir \rightarrow 1024 tables \rightarrow 2^{20} pages

x86-32 Page Mapping: MORE Info

- CR3 \rightarrow 1 page dir \rightarrow 1024 tables \rightarrow 2^{20} pages
 - Up to 1025 pages of metadata per space
 - Just over 4 MiB
- In practice, rarely need all 4 MiB
 - Rare to use entire address space
 - Don't need page tables for unused regions
 - Can reuse page tables for shared mappings (e.g. kernel)
- Most important entry flag: P (Present)
 - Directory entry's P not set? No table
 - Table entry's P not set? No page
 - Try to access not-present page? Page Fault
- Page Fault (#PF, exception 14)

- Handler can load missing page, then resume
- This is swapping



Project Tasks

Tasks/Precode Overview

- You will:
 - Manage virtual memory with dynamic swapping
 - Manage page tables
 - Implement a page fault handler
 - Implement page allocation scheme
 - Swap to/from disk
- Precode infrastructure:
 - Handles saving/restoring CR3 on task switch
 - Installs hooks to fns in `memory.c`:
 - * `void init_memory(void)` — Called on kernel start
 - * `void setup_process_vmem(pcb_t *p)` — Called on process creation
 - * `void page_fault_handler(...)` — Called on page fault
- `memory.c` has been hollowed out
 - Only low-level helper fns remain
 - You must to implement the above hooks and mid-level support fns

Task: Set Up New Process's Address Space

- You may need to...
 - Allocate a physical page
 - Identity map all kernel pages
 - Allocate and "pin" physical pages for page tables if needed
 - Set user access flag for the video buffer page (0xb 8000)
 - Mark all process pages as not present, so they will be loaded on demand
 - Allocate and "pin" a page for user stack
- What do we mean by "pin"?
 - Certain pages are important enough to not swap out
 - You may want to create a "pinned" flag for pages you don't want to swap
- Where do you mark pages as pinned?
 - That's up to you
 - You'll need to track more than the what CPU data structures hold
 - You'll need your own additional data structures

Task: Handle Page Faults (Swap)

- On page fault:
 - CPU stores fault address in CR2

- CPU creates standard exception stack frame (SS, ESP, EFLAGS, CS, EIP, error code)
- Page fault error code contains flags with info about fault
- CPU invokes Page Fault (#PF, exc 14) handler
- Handlers:
 - Low-level handler: `page_fault_handler_entry(...)` [interrupt_asm.S, given]
 - Calls C handler: `page_fault_handler(...)` [memory.c, up to you]
- You may need to...
 - Check error code flags
 - Get fault address from CR2
 - Swap in missing page
 - * Allocate a physical page
 - * Swap out an old page if necessary (write to disk if dirty)
 - * Read needed page from disk
 - * Invalidate page cache for that page

Subtask: Physical Page Allocation

- Remnants of Project 4's page allocation are intact in memory.c
 - You can reuse and enhance them
- If pages are free, simply use a free page and mark it as allocated
- If pages are not free, possible eviction algorithm:


```
loop:
    pick a random page
    if page is not pinned:
        unmap physical page from address space (mark not present)
        if dirty: write page back to disk
        return page for use
```
- Design review Q: What info will you track for allocated pages?
- When swapping out, you can simply write back into the image
 - This will affect the state of the process in the image on reboot
 - Extra challenge: create a separate dedicated swap area

Extra Challenges

- Implement a better page eviction strategy
 - FIFO?
 - FIFO with second chance?
 - Other?
 - See textbook for common algorithms
- Create a dedicated swap area on disk

- Write to swap area instead of overwriting process on disk

Hints

- Be sure to `invalidate_page(...)` [`cpu_x86.h`] when updating mappings
 - Design review Q: When exactly do you need to do this?
- Double-check your bit operations
- Triple-check your protection bits
- Use given helper functions and create your own helper functions
- Use serial I/O for debugging
- Beware: the page fault handler itself can be interrupted by IRQ

Administrative Details

Procedure is (Mostly) the Same as Other Projects

- Design reviews start on Monday (!!)
- Code
- Report
- Hand in via Canvas: zip up entire repository plus report
- **New!-ish** Also upload the report PDF separately

Report

- Structure: scientific paper
- Length: around 4 pages
- Citations: required
 - You must cite sources you use
 - At the very least, you should have the textbook as a source
- File format: PDF
- AI tools: discouraged but not banned. Use must be declared
- **New!-ish** See Howto doc and template in repository
 - `doc/how-to-write-a-report.md`: gives more guidelines and advice
 - `report/latex-src/template.tex`: LaTeX tutorial / template
 - If you use the template, be sure to remove all existing content

Hand In via Canvas

- Put your report in your repository, under a `report/` dir
 - Report should be a PDF format

- If you write in Word or other WYSIWYG word processor, export to PDF
- If you write in a document prep system like Markdown or LaTeX, you can include the report source if you like, but it's not required.
- Zip up your entire repository (code tree + report + .git/ dir)
- Submit via Canvas
 - Upload zip
 - **New!-ish** Also upload report PDF
 - Having both makes it easier for us to grade