



UiT The Arctic University of Norway

Project 5: Virtual Memory

INF-2201 Staff

UiT

Spring 2024

Project 5

Project Overview

Virtual Memory

Project Tasks

Administrative Details

Project 5: Administrative Info (mostly like the others)

- ▶ Mandatory assignment
- ▶ Groups of two
- ▶ Design review
- ▶ Hand-in
 - ▶ Hand-in via Canvas
 - ▶ Code: whole repository (working tree + `.git/` dir)
 - ▶ Report: place in a `report/` directory in your repo
 - ▶ Zip up and upload to Canvas
 - ▶ **New!-ish** Also upload the report PDF separately

Project 5: Virtual Memory

- ▶ In Project 4:
 - ▶ We started enforcing user/kernel separation
 - ▶ Processes had separate address spaces
 - ▶ Virtual memory turned on, but very basic
- ▶ In Project 5:
 - ▶ We will improve the virtual memory
 - ▶ Swap memory pages on demand
- ▶ You will:
 - ▶ Manage virtual memory with dynamic swapping
 - ▶ Manage page tables
 - ▶ Implement a page fault handler
 - ▶ Implement page allocation scheme
 - ▶ Swap to/from disk

Project 5


Project Overview

Virtual Memory

Project Tasks

Administrative Details

Virtual Memory Overview




Mapping from virtual address space
to physical memory

- ▶ Virtual memory
 - ▶ Add indirection!
 - ▶ Map virtual addresses \rightarrow physical memory
- ▶ Divide memory into *pages*
 - ▶ Common size: 4 KiB
 - ▶ 32-bit address space:
 $4 \text{ GiB space} = 4 \text{ KiB/page} \times 2^{20} \text{ pages/space}$
- ▶ Benefits
 - ▶ Use whole address space!
 - ▶ Swap pages to disk as needed!


Multiple Address Spaces

- ▶ Each process can have its own address space
- ▶ Great way to enforce process separation
 - ▶ Can't tamper with what you can't address
 - ▶ Can have per-page permissions
- ▶ But can map some pages into multiple spaces
 - ▶ Common to share kernel pages




Processes with separate address spaces

x86-32 Page Mapping



Page mapping on x86-32

x86-32 Page Mapping: More Info




Page mapping on x86-32

- ▶ To map a 32-bit address space:
 - ▶ 4 GiB space = 4 KiB/page $\times 2^{20}$ pages/space
 - ▶ 20 bits: can select page in space
 - ▶ 12 bits: can select byte in page
- ▶ Two-level hierarchy
 - ▶ Level 2: Page Directory
 - ▶ Level 1: Page Table
 - ▶ Level 0: Physical Page
- ▶ Each step in hierarchy is also a page
 - ▶ 4 KiB page = 32 bits/entry $\times 1024$ entries/page
 - ▶ 20 bits: phys addr of next page in hierarchy
 - ▶ 12 bits: flags
- ▶ CR3 \rightarrow 1 page dir \rightarrow 1024 tables $\rightarrow 2^{20}$ pages

x86-32 Page Mapping: MORE Info

- ▶ CR3 → 1 page dir → 1024 tables → 2^{20} pages
 - ▶ Up to 1025 pages of metadata per space
 - ▶ Just over 4 MiB
- ▶ In practice, rarely need all 4 MiB
 - ▶ Rare to use entire address space
 - ▶ Don't need page tables for unused regions
 - ▶ Can reuse page tables for shared mappings (e.g. kernel)
- ▶ Most important entry flag: P (Present)
 - ▶ Directory entry's P not set? No table
 - ▶ Table entry's P not set? No page
 - ▶ Try to access not-present page? Page Fault
- ▶ Page Fault (#PF, exception 14)
 - ▶ Handler can load missing page, then resume
 - ▶ This is swapping



Page mapping on x86-32

Project 5

Project Overview

Virtual Memory

Project Tasks

Administrative Details

Tasks/Precode Overview

- ▶ You will:
 - ▶ Manage virtual memory with dynamic swapping
 - ▶ Manage page tables
 - ▶ Implement a page fault handler
 - ▶ Implement page allocation scheme
 - ▶ Swap to/from disk
- ▶ Precode infrastructure:
 - ▶ Handles saving/restoring CR3 on task switch
 - ▶ Installs hooks to fns in `memory.c`:
 - ▶ `void init_memory(void)` — Called on kernel start
 - ▶ `void setup_process_vmem(pcb_t *p)` — Called on process creation
 - ▶ `void page_fault_handler(...)` — Called on page fault
- ▶ `memory.c` has been hollowed out
 - ▶ Only low-level helper fns remain
 - ▶ You must to implement the above hooks and mid-level support fns

Task: Set Up New Process's Address Space

- ▶ You may need to...
 - ▶ Allocate a physical page
 - ▶ Identity map all kernel pages
 - ▶ Allocate and “pin” physical pages for page tables if needed
 - ▶ Set user access flag for the video buffer page (0xb 8000)
 - ▶ Mark all process pages as not present, so they will be loaded on demand
 - ▶ Allocate and “pin” a page for user stack
- ▶ What do we mean by “pin”?
 - ▶ Certain pages are important enough to not swap out
 - ▶ You may want to create a “pinned” flag for pages you don't want to swap
- ▶ Where do you mark pages as pinned?
 - ▶ That's up to you
 - ▶ You'll need to track more than the what CPU data structures hold
 - ▶ You'll need your own additional data structures

Task: Handle Page Faults (Swap)

- ▶ On page fault:
 - ▶ CPU stores fault address in CR2
 - ▶ CPU creates standard exception stack frame (SS, ESP, EFLAGS, CS, EIP, error code)
 - ▶ Page fault error code contains flags with info about fault
 - ▶ CPU invokes Page Fault (#PF, exc 14) handler
- ▶ Handlers:
 - ▶ Low-level handler: `page_fault_handler_entry(...)` [interrupt_asm.S, given]
 - ▶ Calls C handler: `page_fault_handler(...)` [memory.c, up to you]
- ▶ You may need to...
 - ▶ Check error code flags
 - ▶ Get fault address from CR2
 - ▶ Swap in missing page
 - ▶ Allocate a physical page
 - ▶ Swap out an old page if necessary (write to disk if dirty)
 - ▶ Read needed page from disk
 - ▶ Invalidate page cache for that page

Subtask: Physical Page Allocation

- ▶ Remnants of Project 4's page allocation are intact in memory.c
 - ▶ You can reuse and enhance them
- ▶ If pages are free, simply use a free page and mark it as allocated
- ▶ If pages are not free, possible eviction algorithm:

loop:

 pick a random page

 if page is not pinned:

 unmap physical page from address space (mark not present)

 if dirty: write page back to disk

 return page for use

- ▶ Design review Q: What info will you track for allocated pages?
- ▶ When swapping out, you can simply write back into the image
 - ▶ This will affect the state of the process in the image on reboot
 - ▶ Extra challenge: create a separate dedicated swap area

Extra Challenges

- ▶ Implement a better page eviction strategy
 - ▶ FIFO?
 - ▶ FIFO with second chance?
 - ▶ Other?
 - ▶ See textbook for common algorithms
- ▶ Create a dedicated swap area on disk
 - ▶ Write to swap area instead of overwriting process on disk

Hints

- ▶ Be sure to `invalidate_page(...)` [`cpu_x86.h`] when updating mappings
 - ▶ Design review Q: When exactly do you need to do this?
- ▶ Double-check your bit operations
- ▶ Triple-check your protection bits
- ▶ Use given helper functions and create your own helper functions
- ▶ Use serial I/O for debugging
- ▶ Beware: the page fault handler itself can be interrupted by IRQ

Project 5

Project Overview

Virtual Memory

Project Tasks

Administrative Details

Procedure is (Mostly) the Same as Other Projects

- ▶ Design reviews start on Monday (!!)
- ▶ Code
- ▶ Report
- ▶ Hand in via Canvas: zip up entire repository plus report
- ▶ **New!-ish** Also upload the report PDF separately

Report

- ▶ Structure: scientific paper
- ▶ Length: around 4 pages
- ▶ Citations: required
 - ▶ You must cite sources you use
 - ▶ At the very least, you should have the textbook as a source
- ▶ File format: PDF
- ▶ AI tools: discouraged but not banned. Use must be declared
- ▶ **New!-ish** See Howto doc and template in repository
 - ▶ `doc/how-to-write-a-report.md`: gives more guidelines and advice
 - ▶ `report/latex-src/template.tex`: LaTeX tutorial / template
 - ▶ If you use the template, be sure to remove all existing content

Hand In via Canvas

- ▶ Put your report in your repository, under a `report/` dir
 - ▶ Report should be a PDF format
 - ▶ If you write in Word or other WYSIWYG word processor, export to PDF
- ▶ Zip up your entire repository (code tree + report + `.git/` dir)
- ▶ Submit via Canvas
 - ▶ Upload zip
 - ▶ **New!-ish** Also upload report PDF
 - ▶ Having both makes it easier for us to grade